
Introduction to HTTP

Topics

HTTP

HTTP/1.x

HTTP/2

HTTP/3

HTTP Request/Response Messages

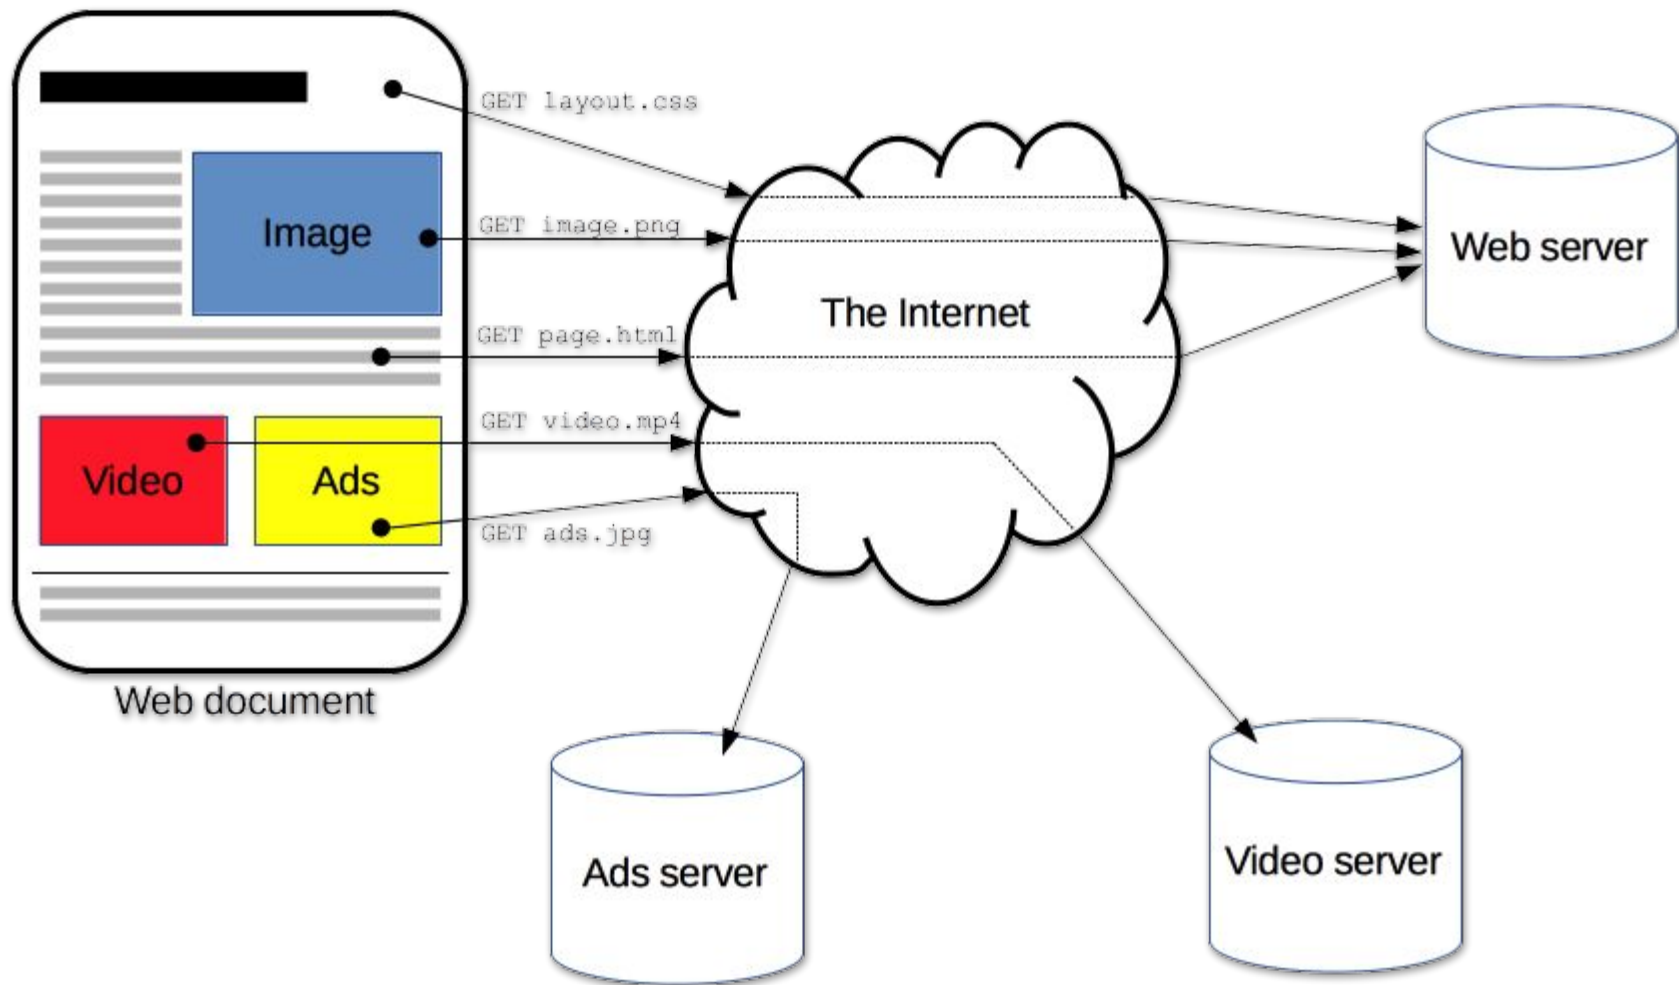
HTTP

A **protocol** for **fetching resources** such as HTML documents

Foundation of any data exchange on the Web

Client-Server protocol, which means requests are initiated by the recipient, usually the Web browser

A complete document is reconstructed from the different sub-documents fetched, for instance, **text, layout description, images, videos, scripts**, and more




HTTP Request/Response

Clients and servers communicate by **exchanging individual messages** (as opposed to a stream of data)

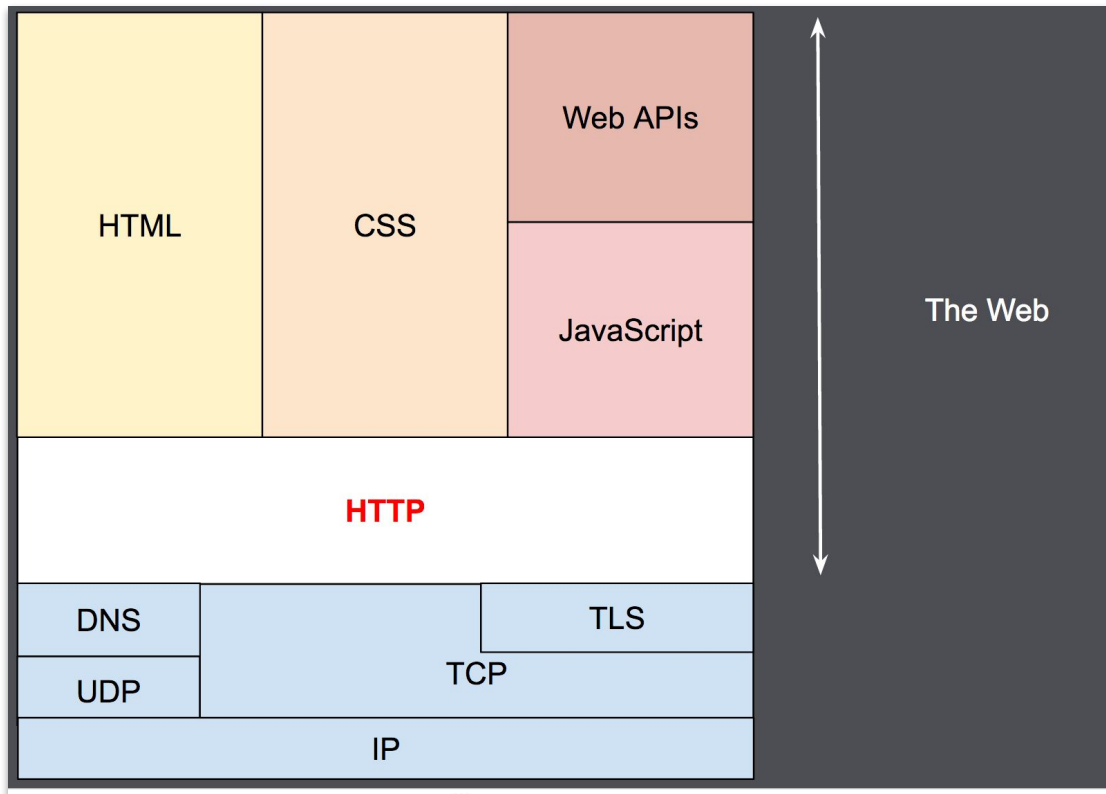
The **messages sent by the client**, usually a Web browser, are called **requests**

The **messages sent by the server** as an answer are called **responses**



HTTP Protocol

HTTP is an application layer protocol that is sent over TCP, or over a TLS-encrypted TCP connection




Components of HTTP-based Systems

Client/User-Agent

any tool that acts on behalf of the user

sends HTTP requests


most of the time the user-agent is a Web browser



Components of HTTP-based Systems

Web Server

handles client requests and provides an answer called the response




Components of HTTP-based Systems

Proxies

sit between the client and the server,

perform functions such as **caching**, **filtering**, **load balancing**, **authentication**, and **logging**



HTTP Statelessness

HTTP is Stateless

there is no link between two requests being successively carried out

However, you can create **sessions** using HTTP Cookies which allows **sharing the same context, or the same state** between requests

HTTP Request/Response Flow

When a client wants to communicate with a server it performs the following steps

Step-1: Open a TCP Connection

The TCP connection is used to send a request(s) and receive a response

The client may open a

new connection,

reuse an existing connection, or

open multiple TCP connections to the servers

HTTP Request/Response Flow

Step-2: Send an HTTP Message

```
GET / HTTP/1.1
```

```
Host: developer.mozilla.org
```

```
Accept-Language: fr
```

HTTP Request/Response Flow

Step-3: Read the Response sent by the Server

```
HTTP/1.1 200 OK
```

```
Date: Sat, 09 Oct 2010 14:28:02 GMT
```

```
Server: Apache
```

```
Last-Modified: Tue, 01 Dec 2009 20:18:22 GMT
```

```
ETag: "51142bc1-7449-479b075b2891b"
```

```
Accept-Ranges: bytes
```

```
Content-Length: 29769
```

```
Content-Type: text/html
```

```
<!DOCTYPE html... (here come the 29769  
bytes of the requested web page)
```

HTTP Request/Response Flow

Step-4: Close or reuse the connection for further requests

Evolution of HTTP

HTTP has gone through many changes

HTTP/0.9

HTTP/1.0

HTTP/1.1

HTTP/2

HTTP/3

HTTP/0.9 – The one-line protocol

Requests consisted of a single line

```
GET /mypage.html
```

The response only consisted of the file itself

```
<html>
```

```
A very simple HTML page
```

```
</html>
```


Characteristics of HTTP/0.9

Client request is a single **ASCII** character string

Client request is terminated by a carriage return (**CRLF**)

Server response is an **ASCII** character stream

Server response is a hypertext markup language (**HTML**)

Connection is terminated after the document transfer is complete

HTTP/0.9

Example

```
$> telnet google.com 80
```

```
Connected to 74.125.xxx.xxx
```

```
GET /about/
```

```
(hypertext response)
```

```
(connection closed)
```

Question

Can you identify some of the limitations of the **HTTP/0.9** protocol ?

HTTP/0.9: Limitations

Could not serve other documents than hypertext documents

It has **GET** request method only

Unable to provide metadata about the request and the response

Unable to negotiate content

HTTP/1.0: Building extensibility

HTTP Working Group (HTTP-WG) published **RFC 1945**, which **documented** the "**common usage**" of the many **HTTP/1.0** implementations found in the internet

HTTP/1.0: Features

Versioning information is now sent within each request

HTTP/1.0 is appended to the **GET** line

The notion of **HTTP headers** has been introduced

both for the requests and the responses, allowing metadata to be transmitted and making the protocol extremely flexible and extensible

HTTP/1.0: Features

Request and response **headers** were **ASCII encoded**

With the help of the new **HTTP headers**, the **ability to transmit other documents than plain HTML files** has been added

using the **Content-Type** header

In addition to media type negotiation it included capabilities such as **content encoding**, **character set support**, **multi-part types**, **authorization**, **caching**, **proxy behaviors**, **date formats**, and more

The connection between server and client is closed after every request

HTTP/1.0 Example

1

Request line with HTTP version number, followed by **request headers**

2

Response status, followed by **response headers**

```
$> telnet website.org 80
```

```
Connected to xxx.xxx.xxx.xxx
```

```
GET /rfc/rfc1945.txt HTTP/1.0 1
```

```
User-Agent: CERN-LineMode/2.15 libwww/2.17b3
```

```
Accept: */*
```

```
HTTP/1.0 200 OK 2
```

```
Content-Type: text/plain
```

```
Content-Length: 137582
```

```
Expires: Thu, 01 Dec 1997 16:00:00 GMT
```

```
Last-Modified: Wed, 1 May 1996 12:45:26 GMT
```

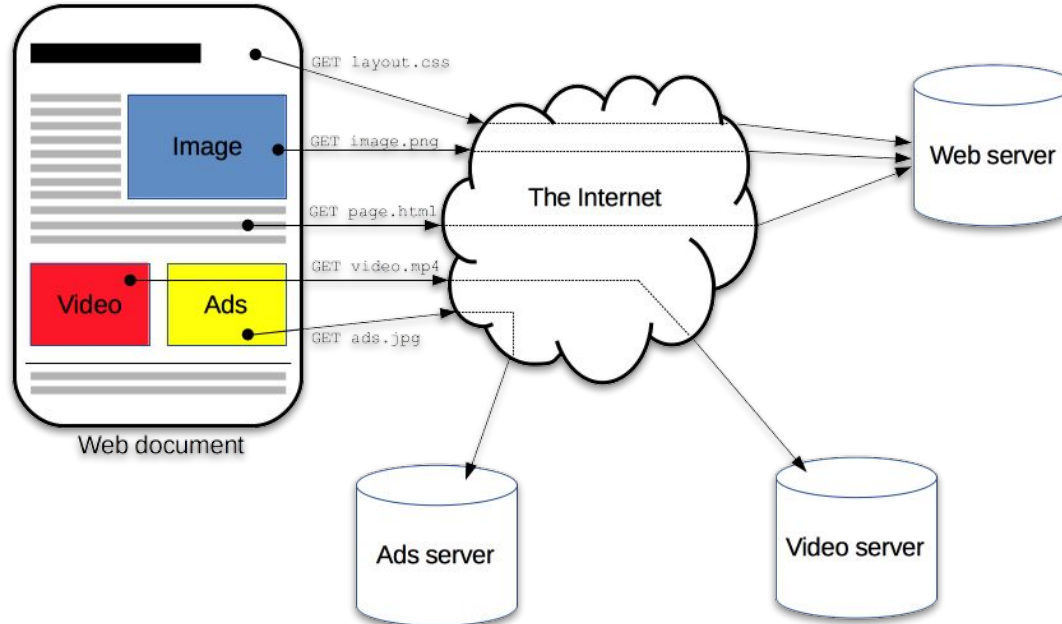
```
Server: Apache 0.84
```

```
(plain-text response)
```

```
(connection closed)
```


Question

What do you think are the **limitations of the HTTP/1.0** protocol ?



HTTP/1.0: Limitations

Requiring a new TCP connection per request

imposes a significant performance penalty

HTTP/1.1: Standardized protocol

The official **HTTP/1.1** standard is defined in **RFC 2616**

The **HTTP/1.1** standard resolved a lot of the protocol ambiguities found in earlier versions

HTTP/1.1 by default leaves the connection open

HTTP/1.1

It introduced a number of critical performance optimizations:

Keepalive Connections

A connection can be reused, saving the time to reopen it numerous times to display the resources embedded into the single original document retrieved

Chunked Encoding Transfers

Chunked responses are now also supported

HTTP/1.1

It introduced a number of critical performance optimizations:

Byte-range Requests

Request Pipelining has been added, allowing to send a second request before the answer for the first one is fully transmitted, lowering the latency of the communication

HTTP/1.1

It introduced a number of critical performance optimizations:

Additional cache control mechanisms have been introduced

Content negotiation, including language, encoding, or type, has been introduced

The addition of the **Host header** allowed to host different domains at the same IP address (allowing server colocation)

HTTP/1.1: Example

- 1 Request for HTML file, with encoding metadata

```
$> telnet website.org 80
```

```
Connected to xxx.xxx.xxx.xxx
```

```
GET /index.html HTTP/1.1 1
```

```
Host: website.org
```

```
User-Agent: Mozilla/5.0 (Macintosh; Intel Mac OS X 10_7_4)... (snip)
```

```
Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8
```

```
Accept-Encoding: gzip,deflate,sdch
```

HTTP/1.1: Example

2 Chunked response for original HTML request

```
HTTP/1.1 200 OK 2
Server: nginx/1.0.11
Connection: keep-alive
Content-Type: text/html; charset=utf-8
Via: HTTP/1.1 GWA
Date: Wed, 25 Jul 2012 20:23:35 GMT
Expires: Wed, 25 Jul 2012 20:23:35 GMT
Cache-Control: max-age=0, no-cache
Transfer-Encoding: chunked
```


HTTP/1.1: Example

- 3 Number of octets in the chunk expressed as an ASCII hexadecimal number
- 4 End of chunked stream response

```
100 3  
<!doctype html>  
(snip)
```

```
100  
(snip)
```

```
0 4
```

HTTP/1.1: Example

5 Request for an icon file made on same TCP connection

6 Inform server that the connection will not be reused

```
GET /favicon.ico HTTP/1.1 5
Host: www.website.org
User-Agent: Mozilla/5.0 (Macintosh; In
Accept: */*
Referer: http://website.org/
Connection: close 6
Accept-Encoding: gzip,deflate,sdch
Accept-Language: en-US,en;q=0.8
Accept-Charset: ISO-8859-1,utf-8;q=0.7
Cookie: __qca=P0-800083390... (snip)
```

HTTP/1.1: Example

7 Icon response, followed by connection close

```
HTTP/1.1 200 OK 7
Server: nginx/1.0.11
Content-Type: image/x-ico
Content-Length: 3638
Connection: close
Last-Modified: Thu, 19 Jul 2012 19:00:00 GMT
Cache-Control: max-age=31536000
Accept-Ranges: bytes
Via: HTTP/1.1 GWA
Date: Sat, 21 Jul 2012 21:00:00 GMT
Expires: Thu, 31 Dec 2037 00:00:00 GMT
Etag: W/PSA-GAu26oXbDi

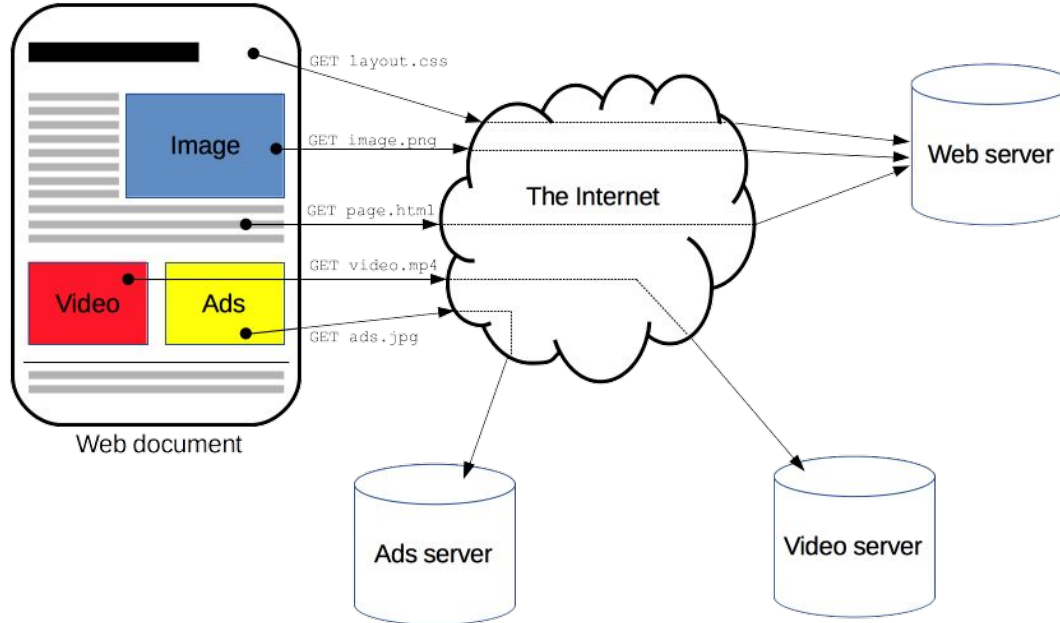
(icon data)
(connection closed)
```

Question

What is the name of the header part that is added on **HTTP/1.1** version to improve the limitation of **HTTP/1.0** (requiring a new TCP connection for each request)

Question

What do you think are the **limitations of the HTTP/1.x** protocol ?



Limitations of HTTP/1.x

Clients need to use **multiple connections to achieve concurrency** and **reduce latency**

Does not compress request and response **headers**, causing unnecessary network traffic

Does not allow effective **resource prioritization**, resulting in poor use of the underlying TCP connection

SPDY

SPDY was an experimental protocol, developed at Google

Its primary goal was to try to **reduce the load latency of web pages** by addressing some of the well-known performance limitations of **HTTP/1.1**

SPDY

The **specific** project **goals** were the following

Target a **50% reduction in page load time (PLT)**

Avoid the need for any changes to content by website authors

Minimize deployment complexity, avoid changes in network infrastructure

Develop this new protocol in partnership with the open-source community

Gather real performance data to (in)validate the experimental protocol

SPDY

SPDY in lab condition has shown **55% reduction in page load time**

As a result SPDY was supported in Chrome, Firefox, and Opera, and a rapidly growing number of sites, both large (e.g., Google, Twitter, Facebook) and small

In effect, SPDY was on track to become a de facto standard through growing industry adoption

SPDY and HTTP/2

Observing the trend, the **HTTP Working Group** (**HTTP-WG**) launched a new effort

- to take the lessons learned from SPDY,
- to build and improve on them, and
- to deliver an official "**HTTP/2**" standard

HTTP/2

HTTP/2 is a protocol designed for **low-latency transport of content** over the World Wide Web


Improve end-user perceived latency

Address the "head of line blocking"


Not require multiple connections

Retain the semantics of HTTP/1.1

Latency vs Bandwidth impact on Page Load Time



Single digit % perf improvement after 5 Mbps



Linear improvement in page load time!

Latency vs Bandwidth impact on Page Load Time

Decreasing latency has more impact than increasing bandwidth

For Example

Decreasing RTTs from 150 ms to 100 ms have a larger effect on the speed of the internet than increasing a user's bandwidth from 3.9 Mbps to 10 Mbps or even 1 Gbps

HTTP/2 : Streams, Messages, and Frames

The introduction of the **new binary framing mechanism** changes how the data is exchanged between the client and server

Stream

A bidirectional flow of bytes within an established connection, which may carry one or more messages

Message

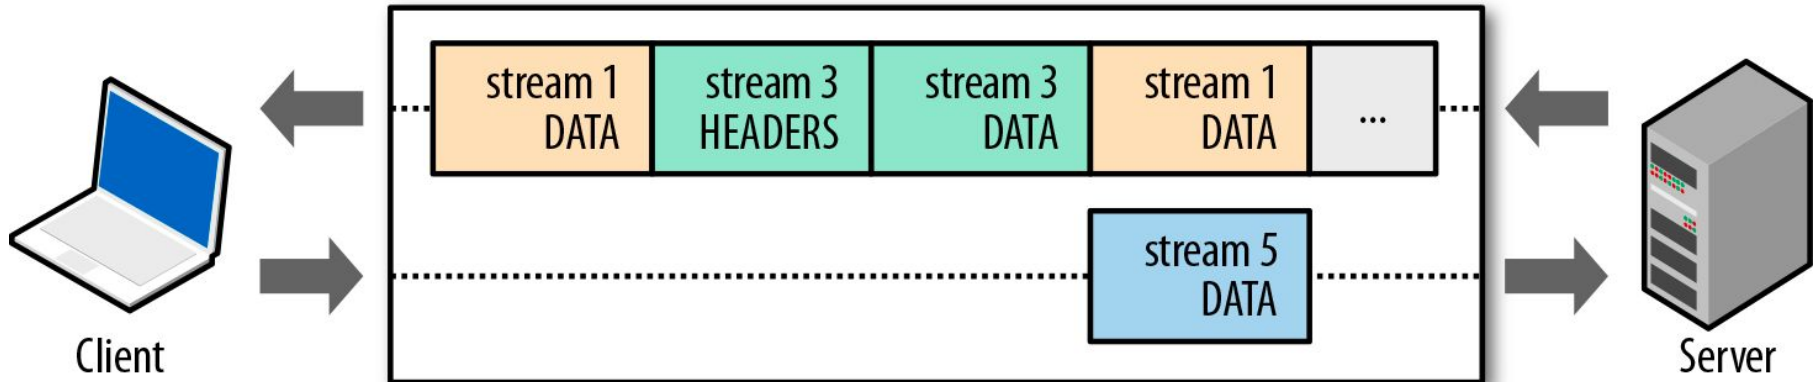
A complete sequence of frames that map to a logical request or response message

HTTP/2: Streams, Messages, and Frames

Frame

The smallest unit of communication in **HTTP/2**, each containing a frame header, which at a minimum identifies the stream to which the frame belongs

HTTP 2.0 connection



HTTP/2 : Streams, Messages, and Frames

The **frame** is the smallest unit of communication that carries a specific type of data—e.g., **HTTP headers**, **message payload**, and so on

Frames from different streams may be **interleaved** and then **reassembled** via the embedded **stream identifier** in the header of each frame

HTTP/2 breaks down the **HTTP protocol communication** into an exchange of **binary-encoded frames**, which are then mapped to **messages** that belong to a particular **stream**, and all of which are **multiplexed** within a **single TCP connection**

HTTP/2 : Main Characteristics

One TCP connection

Stream

Streams are multiplexed

Streams are prioritized

HTTP/2 : Main Characteristics

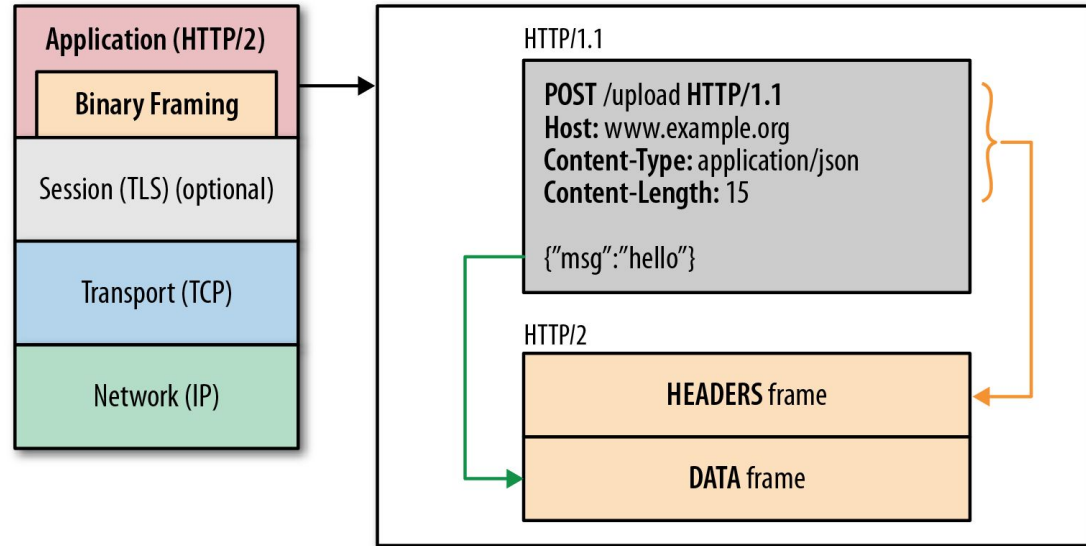
Binary framing layer

Prioritization

Flow control

Server push

Header compression (HPACK)

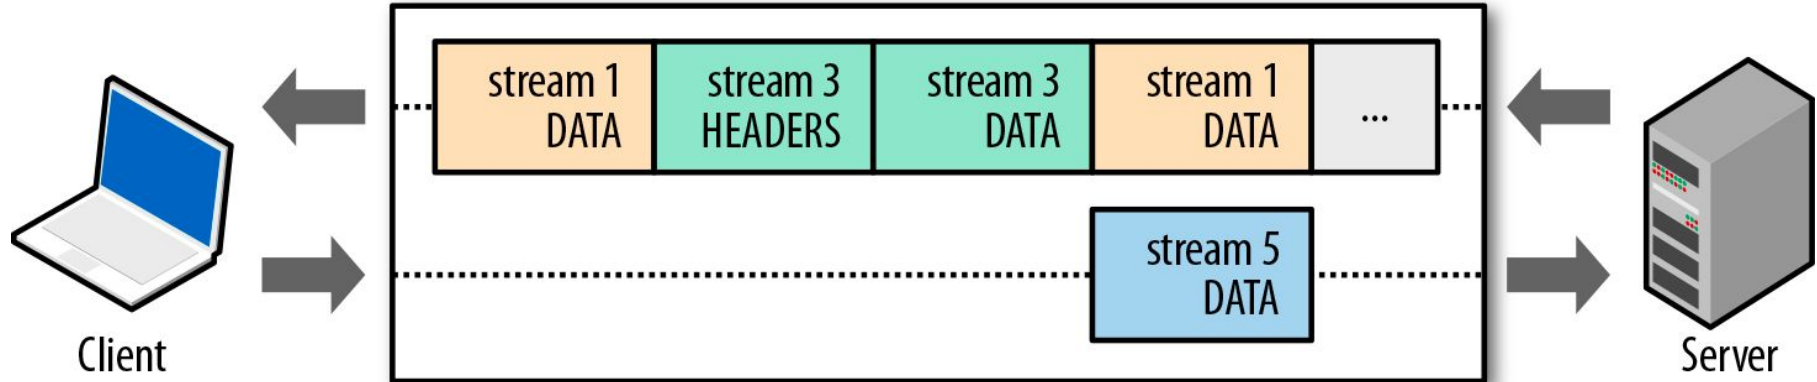


HTTP/2 : Basic data flow

How many streams are there in the diagram?

How many frames?

HTTP 2.0 connection



HTTP/2 : Stream Multiplexing

Advantages

Interleave multiple requests in parallel without blocking on any one

Interleave multiple responses in parallel without blocking on any one

Use a **single connection to deliver multiple requests** and responses in parallel

HTTP/2 : Stream Multiplexing

Advantages

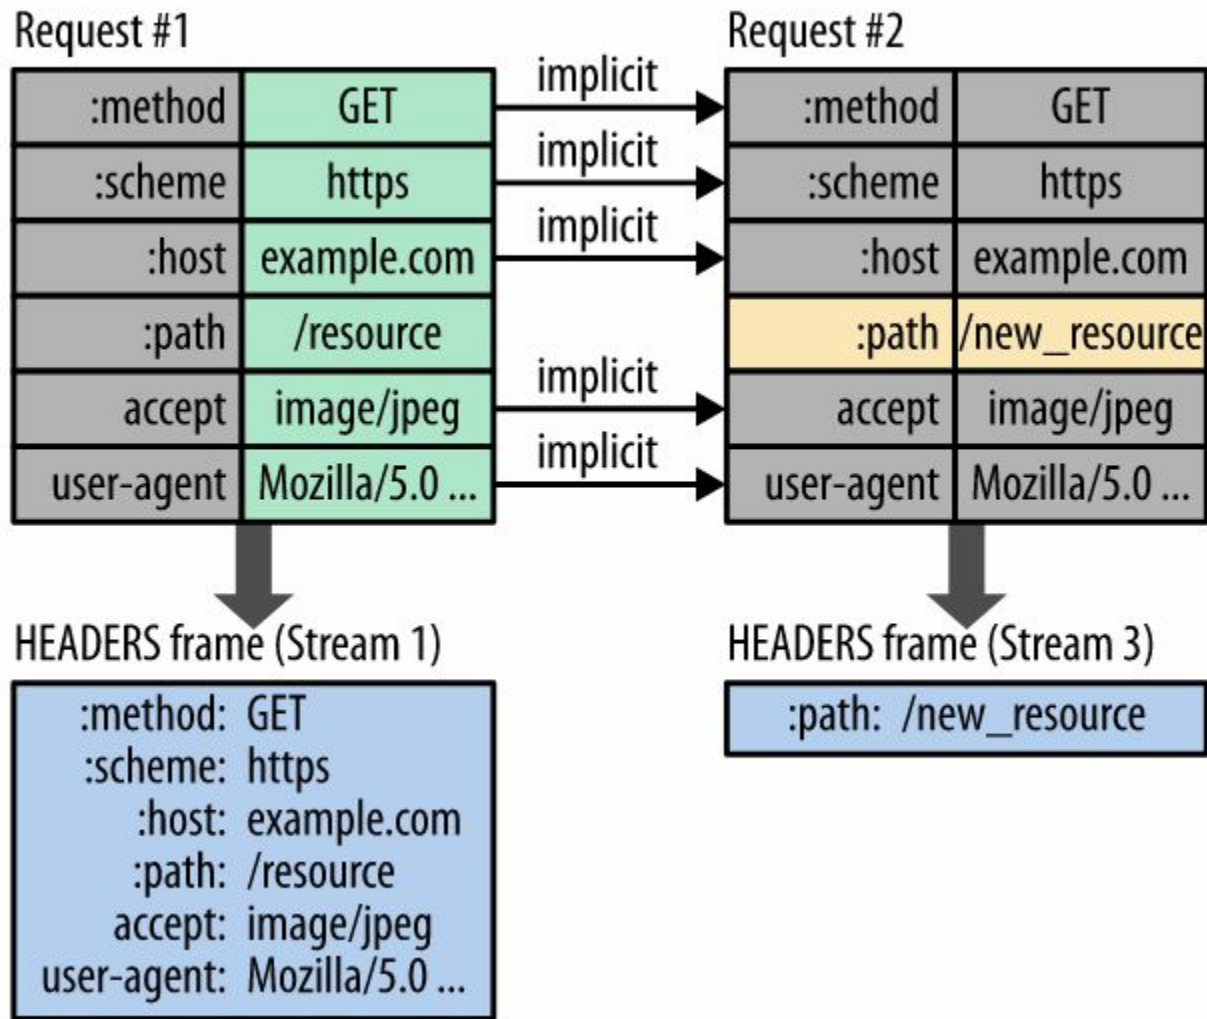
Remove unnecessary HTTP/1.x workarounds for optimization, such as concatenated files, image sprites, and domain sharding

Deliver **lower page load times** by eliminating unnecessary latency and improving utilization of available network capacity

HTTP/2

Header Compression

Uses **HPACK** algorithm

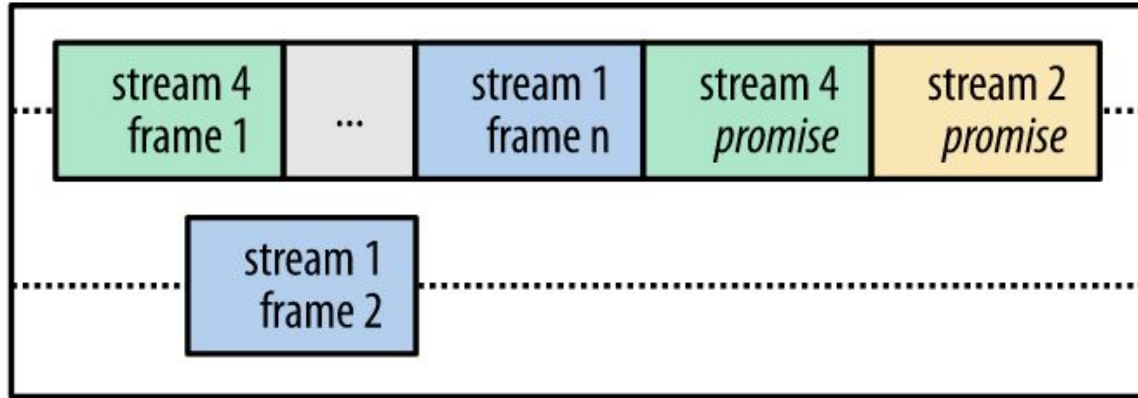


HTTP/2 : Server Push

HTTP 2.0 connection



stream 1
frame 1



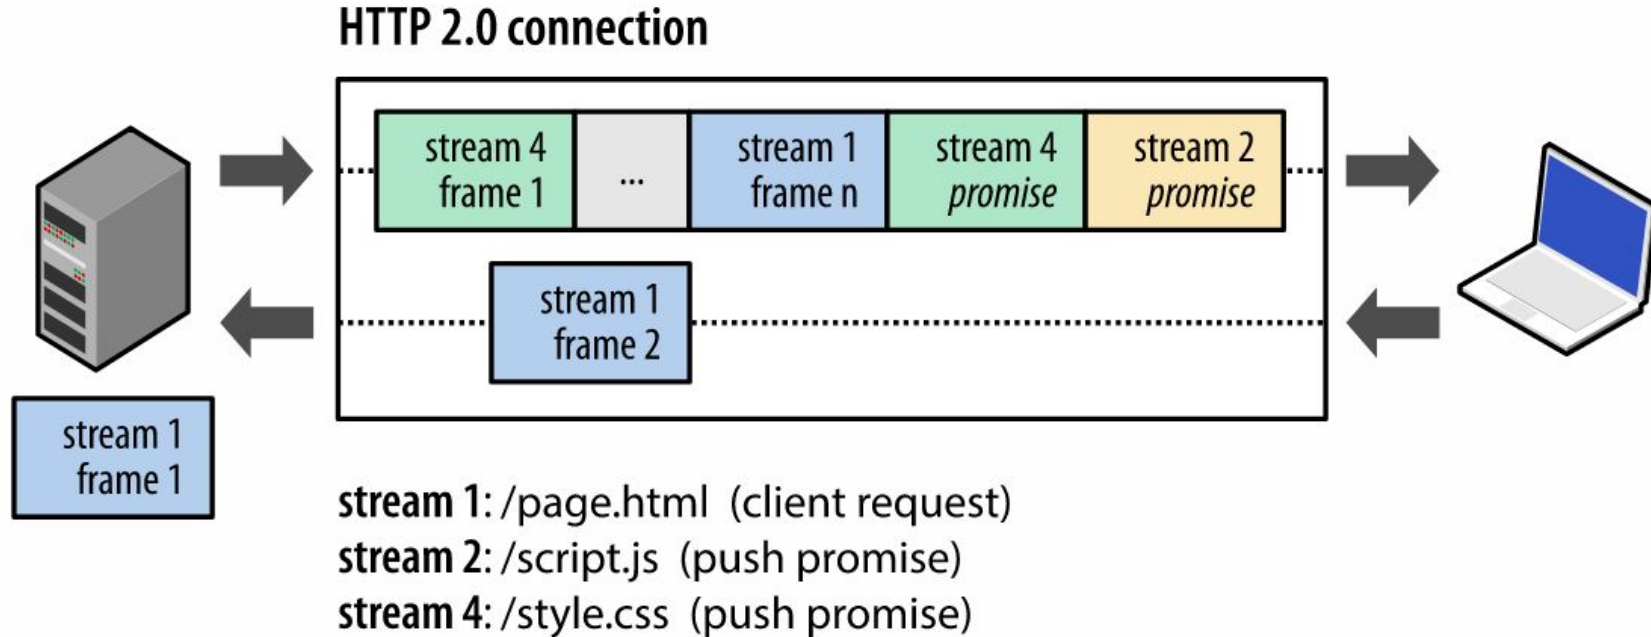
stream 1: /page.html (client request)

stream 2: /script.js (push promise)

stream 4: /style.css (push promise)

HTTP/2 : Server Push

What are the advantages of server push?



HTTP/2 : Advantages of Server Push

Pushed resources can be cached by the client

Pushed resources can be reused across different pages

Pushed resources can be multiplexed alongside other resources

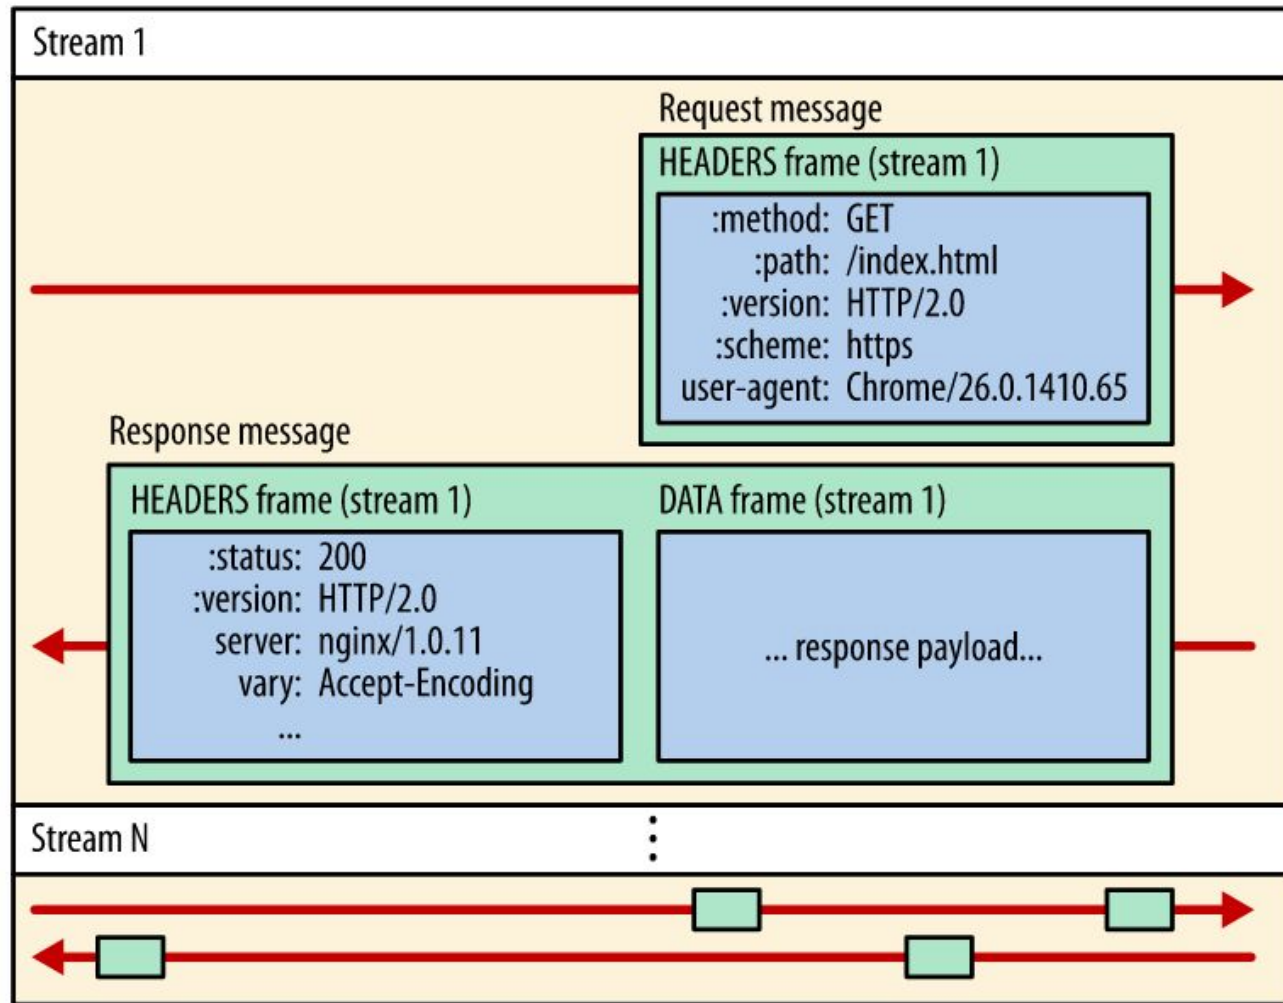
Pushed resources can be prioritized by the server

HTTP/2

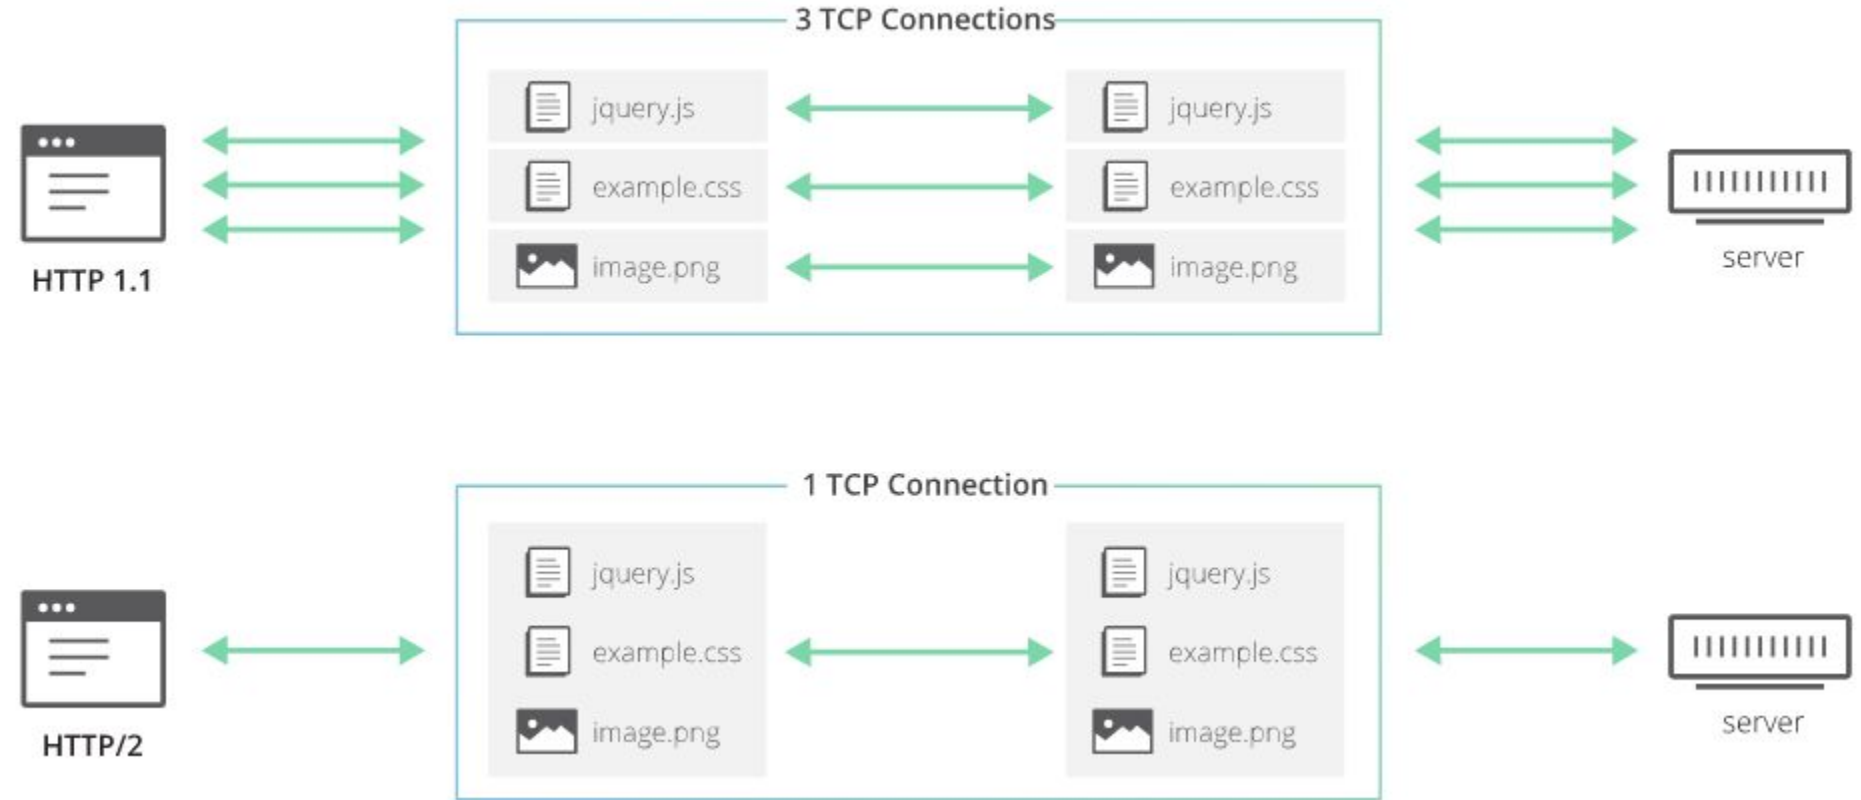
Single TCP

Multiple Stream

Connection



HTTP/1.X vs HTTP/2 TCP Connection



Limitations of HTTP/2

Can you think of any limitation of HTTP/2 protocol?

QUIC

QUIC (Quick **UDP** Internet **C**onnections) is a new transport protocol for the internet, developed by Google

QUIC solves a number of transport-layer and application-layer problems experienced by modern web applications, while requiring little or no change from application writers

QUIC is very similar to **TCP+TLS+HTTP2**, but **implemented on top of UDP**

QUIC

Key **advantages** of QUIC over **TCP+TLS+HTTP2** include:

- Low connection establishment latency

- Improved congestion control

- Multiplexing without head-of-line blocking

- Forward error correction

- Connection migration

HTTP / 3: HTTP over QUIC

Instead of using TCP as the transport layer for the session, it uses QUIC

QUIC introduces streams as first-class citizens at the transport layer

QUIC streams share the same QUIC connection, so no additional handshakes and slow starts are required to create new ones

QUIC streams are delivered independently such that in most cases packet loss affecting one stream doesn't affect others.

This is possible because QUIC packets are encapsulated on top of UDP datagrams

HTTP/3: HTTP over QUIC

Using **UDP** allows **much more flexibility** compared to **TCP**, and enables **QUIC** implementations to live fully in user-space — updates to the protocol implementations are not tied to operating systems updates as is the case with TCP

QUIC also combines the typical 3-way TCP handshake with **TLS 1.3**'s handshake

Encryption and **authentication** are provided by default, and also enables **faster connection establishment**

HTTP/3: HTTP over QUIC

Check the browsers supporting **QUIC** or **HTTP/3**

[CanIUse](#)

HTTP Messages

HTTP messages, as defined in HTTP/1.1 and earlier, are human-readable

In **HTTP/2**, these messages are embedded into a binary structure, a frame, allowing optimizations like compression of headers and multiplexing

Even if only part of the original HTTP message is sent in **HTTP/2**, **the semantics of each message is unchanged** and the client reconstitutes the original HTTP/1.1 request

HTTP messages typically contain, **request/response line**, **request/response headers**, and/or **request/response body**

HTTP Messages

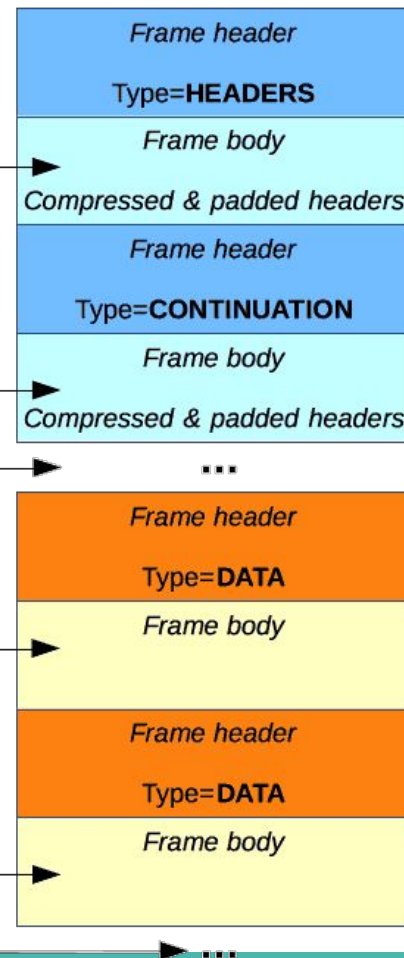
HTTP/1.X vs HTTP/2

HTTP/1.x message

```
PUT /create_page HTTP/1.1
Host: localhost:8000
Connection: keep-alive
Upgrade-Insecure-Requests: 1
Content-Type: text/html
Content-Length: 345
```

```
Body line 1
Body line 2
...
```

HTTP/2 stream (composed of frames)



HTTP Request Message

```
GET /doc/test.html HTTP/1.1
```

```
Host: www.test101.com
```

```
Accept: image/gif, image/jpeg, */*
```

```
Accept-Language: en-us
```

```
Accept-Encoding: gzip, deflate
```

```
User-Agent: Mozilla/4.0
```

```
Content-Length: 35
```

```
bookId=12345&author=Tan+Ah+Teck
```

Request Line

Request Headers

Request
Message
Header

A blank line separates header & body

Request Message Body

HTTP Response Message

HTTP/1.1 200 OK

Date: Sun, 08 Feb xxxx 01:11:12 GMT

Server: Apache/1.3.29 (Win32)

Last-Modified: Sat, 07 Feb xxxx

ETag: "0-23-4024c3a5"

Accept-Ranges: bytes

Content-Length: 35

Connection: close

Content-Type: text/html

<h1>My Home page</h1>

Status Line

Response Headers

Response
Message
Header

A blank line separates header & body

Response Message Body

Reference

High Performance Browser Networking

HTTP | MDN

<https://www.rfc-editor.org/rfc/rfc9114.html>

<https://www.rfc-editor.org/rfc/rfc9000.html>

Introduction to Server Side Programming

Topics

Server-Side/Backend Programming


Static/Dynamic Websites

Server-Side Web Frameworks



Users

Collect Data




Display Results



What user sees
& interacts with
HTML, CSS, JavaScript

Frontend

Request



Response



Contains App Logic
PHP, JavaScript, Python, Java

Web Server



File System

HTML, CSS, Images



Database

MySQL, PostgreSQL,
MariaDB

Backend

Web Application Architecture

Server Side Programming

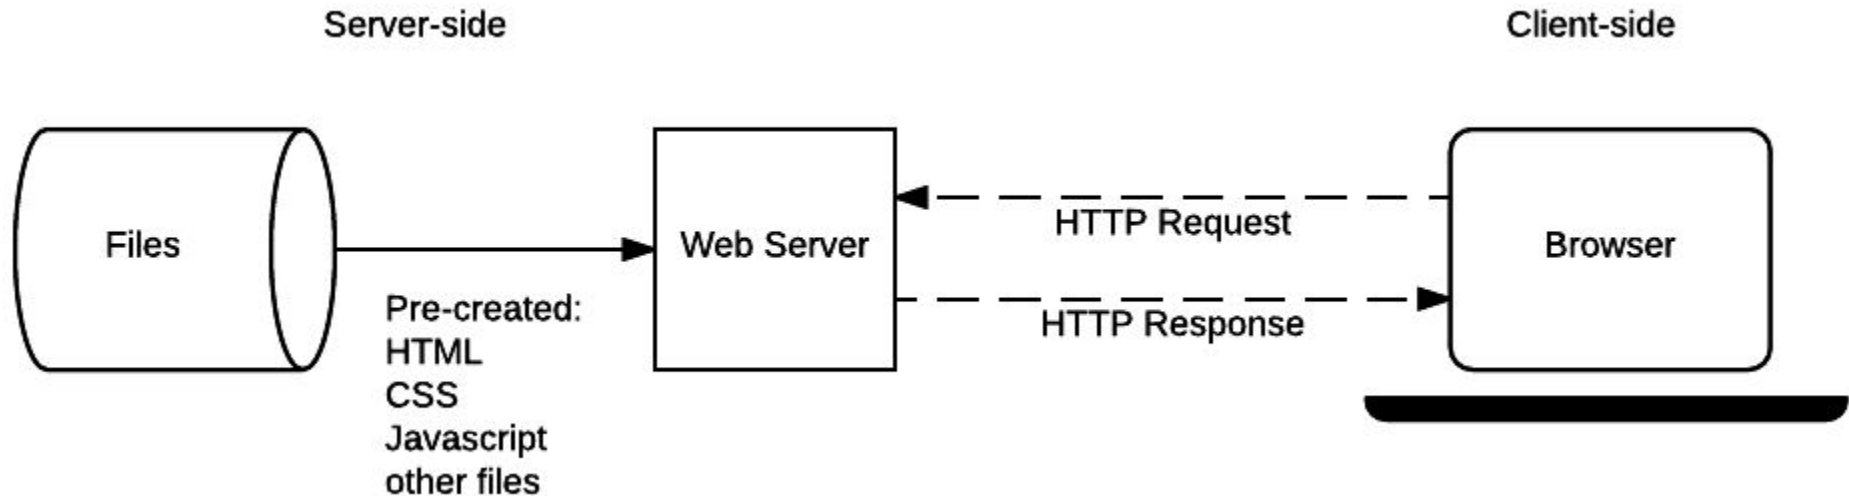
Dynamically display different data when needed, generally pulled out of a database stored on a server and sent to the client to be displayed via some code (e.g. HTML and JavaScript)

Web Servers

Web servers wait for client request messages, process them when they arrive, and reply to the web browser with an HTTP response message

Static Sites

Returns the same hard-coded content from the server whenever a particular resource is requested



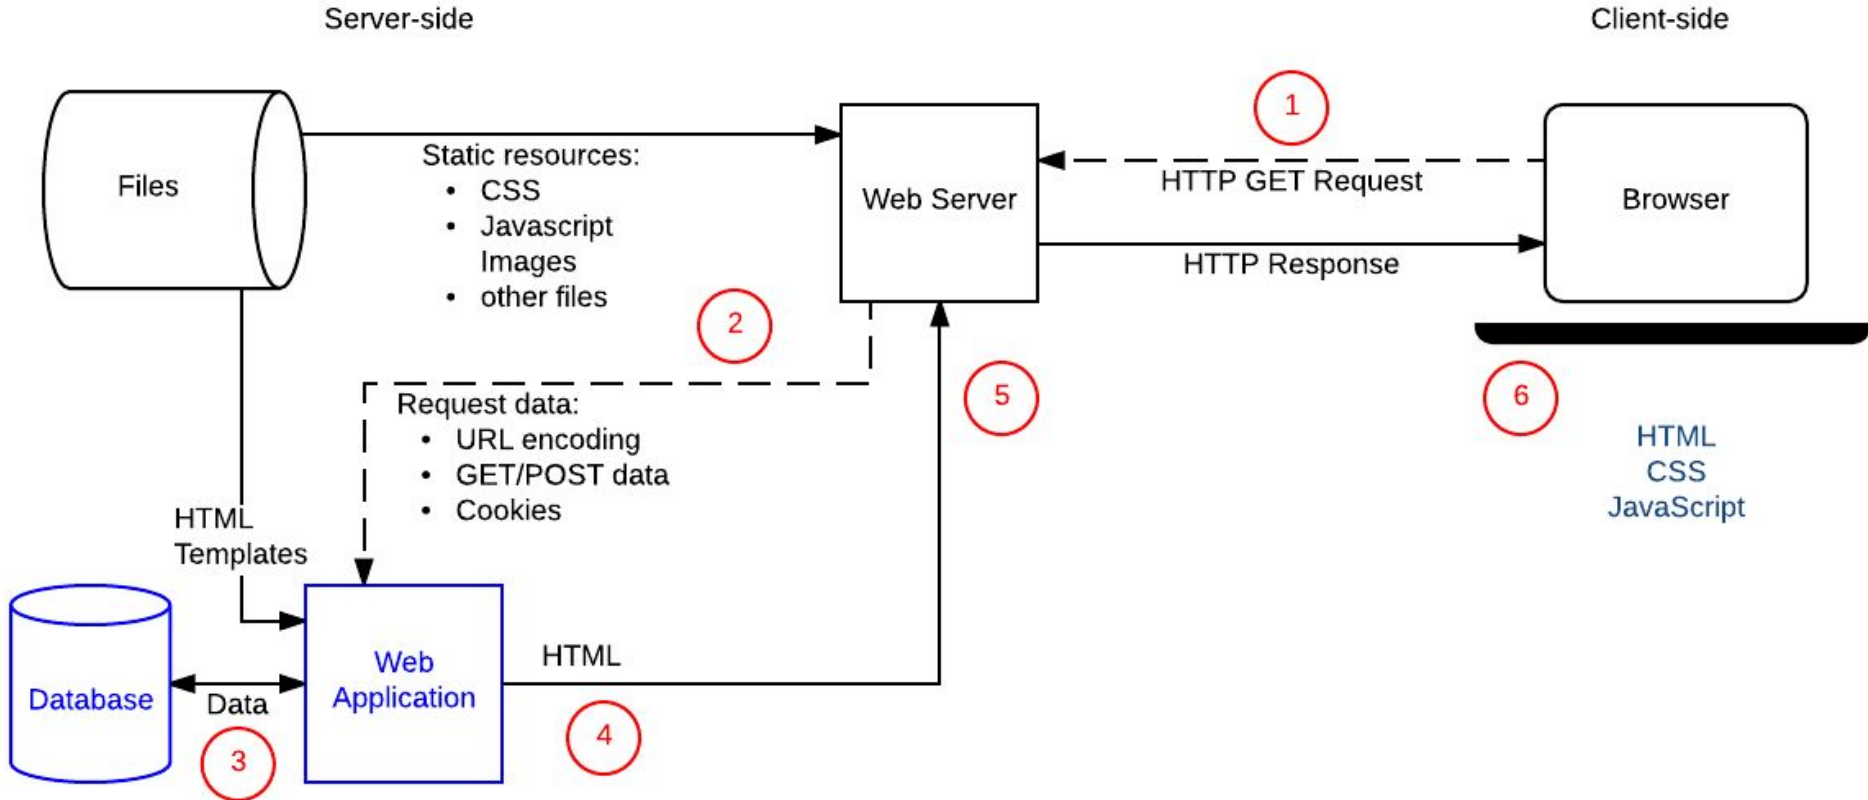
Dynamic Sites

Some of the **response** content is **generated dynamically**, only when needed

HTML pages are normally created by inserting data from a database into placeholders in **HTML templates**

Can return different data for a URL based on information provided by the user or stored preferences and can perform other operations as part of returning a response

Dynamic Sites



What can you do on the server-side?

Efficient Storage and Delivery of Information

Customised User Experience

servers can store and use information about clients to provide a convenient and tailored user experience

Controlled Access to Content

restrict access to authorized users and serve only the information that a user is permitted to see

What can you do on the server-side?

Store Session/State Information

server-side programming allows developers to make use of sessions

session is a mechanism that allows a server to store information on the current user of a site and send different responses based on that information

What can you do on the server-side?

Notifications and Communication

servers can send general or user-specific notifications through the website itself or via email, SMS, instant messaging, video conversations, or other communications services

What can you do on the server-side?

Data Analysis

A website may collect a lot of data about users: what they search for, what they buy, what they recommend, how long they stay on each page

Server-side programming can be used to refine responses based on analysis of this data

Server-Side Web Frameworks

Also known as web application frameworks

They make it easier to write, maintain and scale web applications

Server-Side Web Frameworks

Provide tools and libraries that simplify common web development tasks, such as

- route requests** to the appropriate handler

- make it easy to **access data in the request**

- abstract and simplify **database access**

- formatting** output (e.g. HTML, JSON, XML)

- improving **security** against web attacks

Some Web Frameworks

Framework	Language and Environment
<u>Django</u>	Python
<u>Flask</u>	Python
<u>Express.js</u>	NodeJs, Javascript
<u>Nest.js</u>	NodeJs, Express.js/Fastify, Javascript
<u>Deno</u>	JavaScript, TypeScript, Chrome v8, Rust
<u>Ruby on Rails</u>	Ruby

Some Web Frameworks

Framework	Language and Environment
<u>ASP.NET</u>	.NET, C#
<u>Micronaut</u>	Java
<u>Quarkus</u>	Java
<u>Spring Boot</u>	Java, Kotlin
<u>Laravel</u>	PHP

Reference

https://developer.mozilla.org/en-US/docs/Learn/Server-side/First_steps/Introduction